Gmail

benjamin barber <starworks5@gmail.com>

Deontic Logic and Legal systems

8 messages

benjamin barber <starworks5@gmail.com> Thu, Oct 16, 2025 at 2:01 PM
To: bcahan@stanford.edu

https://www.cambridge.org/core/books/deontic-logic-and-legal-systems/AD6FBAF0E92AAA3A241403284FE4E73A
https://github.com/endomorphosis/ipfs_datasets_py?tab=readme-ov-file#%EF%B8%8F-legal-document-formalization

image.png
388K

Bruce Cahan <bcahan@stanford.edu> Thu, Oct 16, 2025 at 2:10 PM
To: benjamin barber <starworks5@gmail.com>

Hi Ben:

Thanks for sharing this work.

Can we set up a Zoom to map synergies?

All the best,

Bruce

Bruce Cahan
Cell: (212) 399-9700
bcahan@urbanlogic.org | bcahan@stanford.edu
Skype: bbcesq
Twitter: brucecahan

Ashoka Fellow - Social Investment Entrepreneur
Stanford University Department of Management Science & Engineering - Lecturer
President - Urban Logic ® (A NY nonprofit, qualified in CA) | www.urbanlogic.org
Profile - https://www.linkedin.com/in/brucecahan/
Curator/Convener –www.TEDxNewWallStreet.org
ORCID: https://orcid.org/0000-0001-6903-9403

> On Oct 16, 2025, at 2:01 PM, benjamin barber <starworks5@gmail.com> wrote:
>
> https://www.cambridge.org/core/books/deontic-logic-and-legal-systems/AD6FBAF0E92AAA3A241403284FE4E73A
> https://github.com/endomorphosis/ipfs_datasets_py?tab=readme-ov-file#%EF%B8%8F-legal-document-

formalization

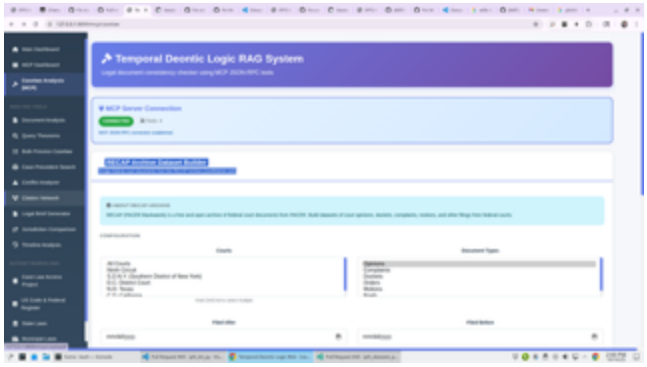

---

benjamin barber <starworks5@gmail.com> Thu, Oct 16, 2025 at 2:21 PM
To: Bruce Cahan <bcahan@stanford.edu>

Just let me know what your calendar looks like, or if it's easier for you, we can meet at Yannic Kilcher's (CEO of Deep Judge) Weekend Machine learning peer review paper discussion on discord on Saturday 11 am PST. https://discord.gg/Ub6cx5sU

There are normally 100 - 150 people in attendance on the weekends, to review papers from the machine learning preprint Arxiv, most of them are domain experts, and a few are chairs of Neurips, retired CERN / Sandia physicists, etc, but Myself (hallucinate llc) and Yannic (deepjudge) and Xmaster (Aleph Alpha) work in the legal domain.

11:16 5G

Provable Scaling Laws of Feature

Saturday Paper Discussion 1 118

LIVE

such as Muon can be effective, from the first principles of gradient dynamics. Our analysis can be extended to multi-layer architectures.

1 INTRODUCTION

While modern deep models such as Transformers have achieved impressive empirical performance, it remains a mystery how such models acquire the knowledge during the training process. There have been ongoing arguments on whether the models can truly generalize beyond what it is trained on, or just memorize the dataset and performs poorly in out-of-distribution (OOD) data (Wang et al., 2024b; Chu et al., 2025; Mirzadeh et al., 2024).

Modeling the memorization/generalization behaviors have been a goal of many works. One such behavior, know as *grokking* (Power et al., 2022; Doshi et al., 2024; Nanda et al., 2023; Wang et al., 2024a; Varma et al., 2023; Liu et al., 2023; Thilak et al., 2022), shows that the model initially overfits to the training set, and then suddenly generalizes to unseen test samples after continuous training. Many explanation exists, e.g., effective theory (Liu et al., 2022; Clauw et al., 2024), efficiency of memorization and generalization circuits (Varma et al., 2023), Bayesian interpretation with weight decay as prior (Millidge, 2022), etc. Most works focus on a direct explanation of its empirical behaviors, or leveraging property of very wide networks (Barak et al., 2022; Mohamadi et al., 2024; Rubin et al., 2024), but few explores the details of the grokking learning procedure by studying the gradient dynamics on the weights.

In this [illegible] mathematical framework $\mathtt{Li}_2$ that divides the grokking dynamics for 2-layer nonlinear networks into three major stages (Fig. 1). *Stage I: Lazy Learning*: when training

Yannic Kilcher

Speakers — 1

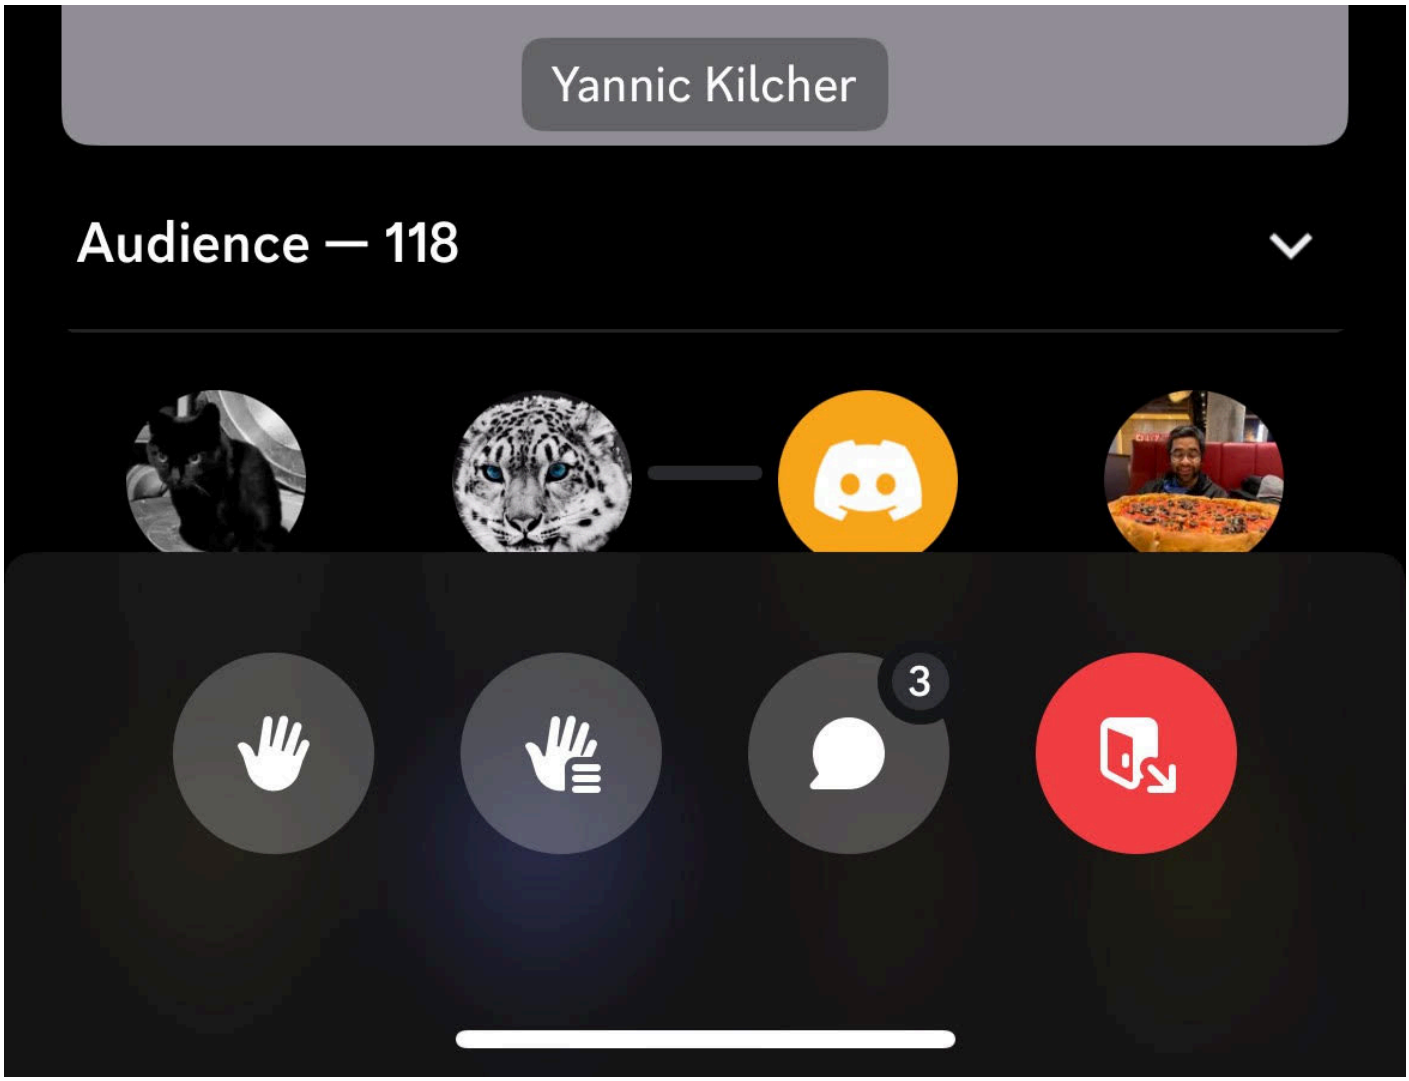
Yannic Kilcher
Audience — 118
3

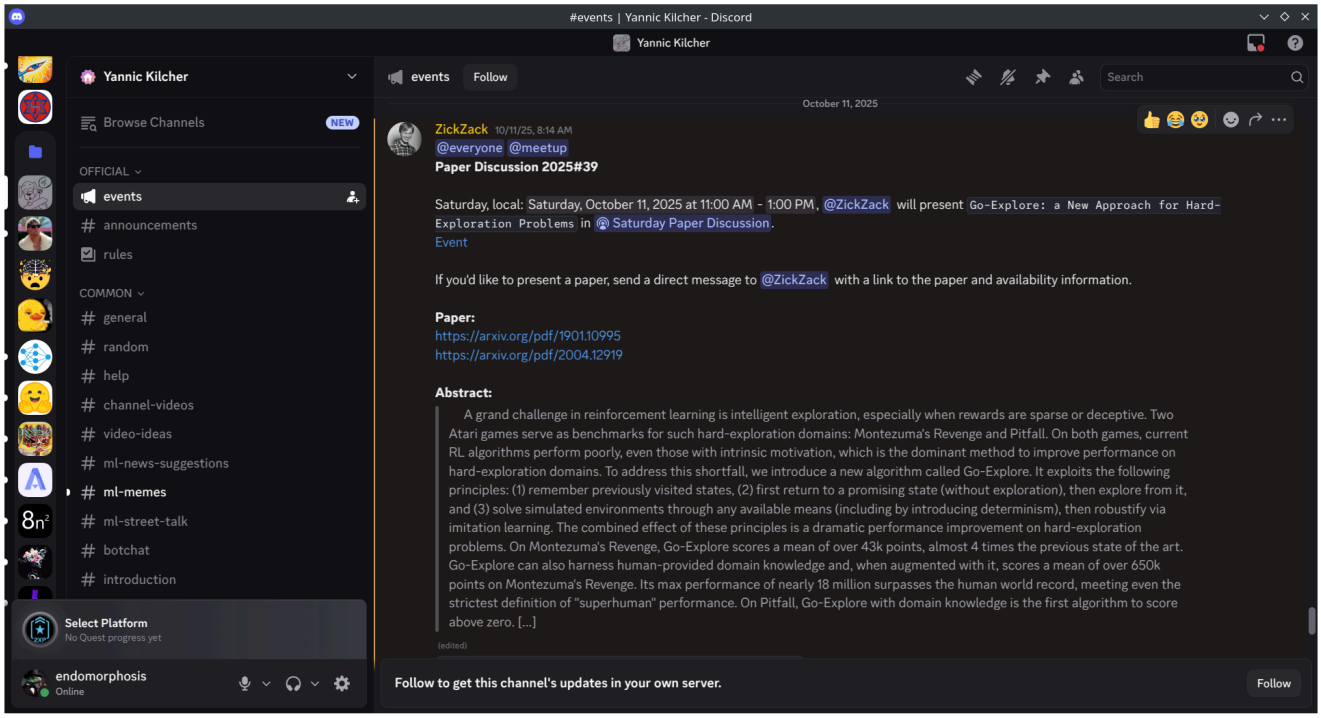

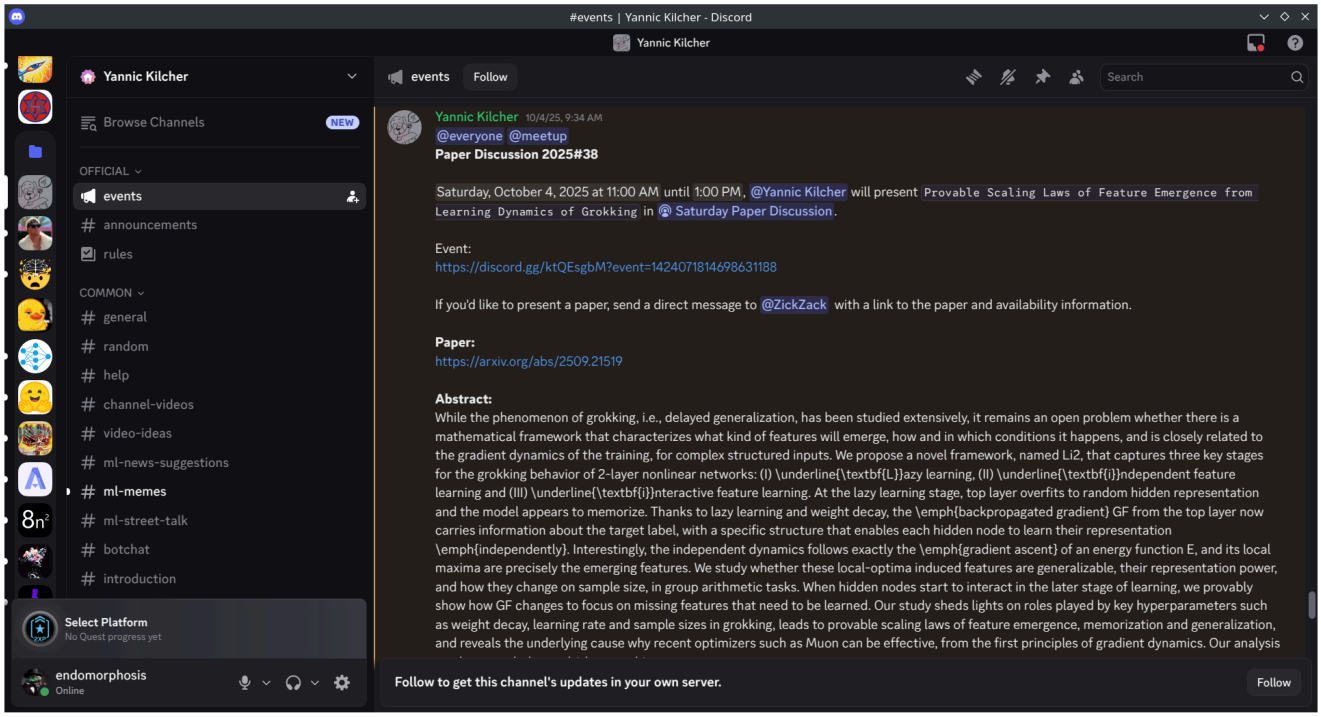

[Quoted text hidden]

---

benjamin barber <starworks5@gmail.com> Thu, Oct 16, 2025 at 7:59 PM
To: Bruce Cahan <bcahan@stanford.edu>

Since apparently you're into engineering.

On Thu, Oct 16, 2025 at 2:10 PM Bruce Cahan <bcahan@stanford.edu> wrote:
[Quoted text hidden]

Complex_Logical_Query_Answering.pdf
14966K

---

benjamin barber <starworks5@gmail.com> Thu, Oct 16, 2025 at 8:06 PM
To: Bruce Cahan <bcahan@stanford.edu>

Portland state university mentorship

https://braindump.cat.pdx.edu/

You can also find me at Linux foundation infiniedge ai meetings

https://lfedge.org/projects/infiniedge-ai/

And with LAION.ai

https://www.intel.com/content/www/us/en/developer/articles/technical/bud-e-ai-assisted-education-for-all.html

[Quoted text hidden]

5 attachments

IMG_0521.jpeg
1204K

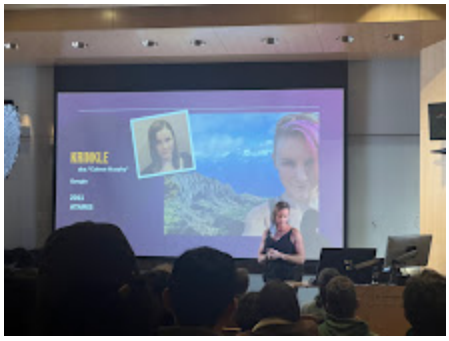

IMG_0522.jpeg
1295K

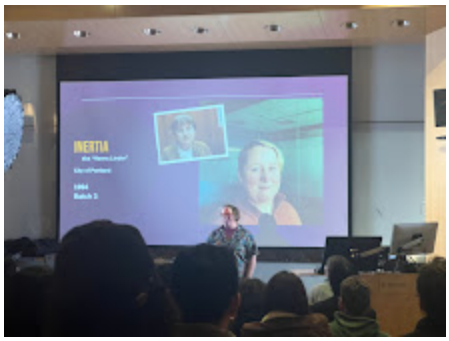

IMG_0523.jpeg
1404K

IMG_0519.jpeg
2899K

IMG_0520.jpeg
1589K

---

benjamin barber <starworks5@gmail.com> Thu, Oct 16, 2025 at 8:09 PM
To: Bruce Cahan <bcahan@stanford.edu>

https://events.linuxfoundation.org/open-source-ai-week/#infiniedge-ai

InfiniEdge AI @ Open Source AI Week 2025
was updated

OCT 23 **Thursday, October 23**
4:30 PM - 6:30 PM PDT

3120 Scott Blvd ↗
Santa Clara, California

The following information has been updated:

- The end time changed from Thursday, Oct 23, 6:00 PM PDT to **Thursday, Oct 23, 6:30 PM PDT**.

View Event

We've attached a new calendar event which should automatically update your calendar.

[Quoted text hidden]

---

benjamin barber <starworks5@gmail.com> Thu, Oct 16, 2025 at 8:35 PM
To: Bruce Cahan <bcahan@stanford.edu>

https://www.youtube.com/watch?v=JIVKiMLJ5bE

A public defender office in Marion County is suing the Oregon agency in charge of indigent defense, claiming their new contract would force attorneys to behave unethically by taking on excessive caseloads. (3000 unrepresented defendants)
https://www.kgw.com/article/news/local/the-story/oregon-public-defender-marion-county-lawsuit-state-opdc-contract-caseload/283-31ef3b65-78b5-40b0-974b-6c9c148fb9fb

[Quoted text hidden]

---

benjamin barber <starworks5@gmail.com> Sat, Oct 18, 2025 at 9:34 AM
To: Bruce Cahan <bcahan@stanford.edu>

9:32

events ›

ZickZack 9:31 AM

@everyone @meetup

Paper Discussion 2025#40

Saturday, local: Saturday, October 18, 2025 at 11:00 AM - 1:00 PM, **@wave function** will present `Towards Fast, Accurate Value Estimates for CoT RL` in **Saturday Paper Discussion**.
Event

If you'd like to present a paper, send a direct message to **@ZickZack** with a link to the paper and availability information.

Paper:
research in development

Abstract:

> The VinePPO paper [1] investigated the

standard value estimations used in Reinforcement Learning with Verifiable Rewards (RLVR) for Chain-of-Thought (CoT) reasoning, and found that they produce poor estimates, barely outperforming random baselines when ranking alternative steps, and impairing training and final reasoning

Follow to get this channel's updates in your own server.

Follow

@everyone <@&782362139087208478>
Paper Discussion 2025#40

Saturday, local: <t:1760810400:F> - <t:1760817600:t>, <@294632311176298498> will present `Towards Fast, Accurate Value Estimates for CoT RL` in <#1337720641528401930>.
[Event](https://discord.com/events/714501525455634453/1429143920817995896)

If you'd like to present a paper, send a direct message to <@178150411545346049> with a link to the paper and availability information.

Paper:
research in development

Abstract:
> The VinePPO paper [1] investigated the standard value estimations used in Reinforcement Learning with Verifiable Rewards (RLVR) for Chain-of-Thought (CoT) reasoning, and found that they produce poor estimates, barely outperforming random baselines when ranking alternative steps, and impairing training and final reasoning performance. We study whether, for the PPO value-network model in the VinePPO paper's baseline, this is a training-data problem or a fundamental architectural limitation, by using VinePPO's accurate Monte-Carlo (MC) estimates as supervision targets for the value network. Ideally, we would like to develop a hybrid training approach for even more efficient PPO-like RL, where the value network is regressed against the training outcomes most of the time, as usual, and only occasionally regressed against the VinePPO MC rollouts

[Quoted text hidden]