ChatGPT Log in Sign up for free

what do you think about the skill of the github user "endomorphosis"

Based on what's publicly visible, the GitHub user **endomorphosis** (Benjamin J Barber) looks like a **highly active, experimental engineer in the AI + decentralized systems space**—but with a very particular style that's worth unpacking.

🧠 What they're good at

1. Ambitious systems thinking

A consistent theme across their work is combining:

- **AI/ML (Hugging Face, LLMs)**
- **Distributed systems (IPFS, libp2p)**
- **Inference orchestration across nodes**

For example, projects like:

- IPFS-based model serving + sharing
- Distributed inference fallback to cloud/docker providers

show **strong architectural thinking**, not just coding. PyPI

That's not beginner-level—this is **systems design territory**.

---

2. Breadth of technologies

They list experience across:

- Python, Node.js, C++, Java, etc. Devpost - T...

And their repos span:

- APIs (FastAPI servers)
- ML tooling
- Networking layers
- Data pipelines

That suggests a **generalist who can build full stacks**, especially in infra-heavy ML.

---

3. Output volume and experimentation

They've:

- Published multiple libraries and bridges (transformers, datasets, FAISS, etc.)
- Built custom tooling around OpenVINO and distributed inference
- Uploaded datasets and models on Hugging Face Hugging Face

This level of output usually indicates:

- Curiosity
- Willingness to prototype fast
- Comfort working without strict structure

⚠️ Where things are weaker (or at least unclear)

1. Polish and adoption

There's little evidence of:

- Large community adoption
- High GitHub stars / production usage
- Widely recognized libraries

That often means the work is:

- **Experimental or early-stage**
- Not heavily refined for production or usability

---

2. Scope vs. execution risk

A lot of the ideas are **very ambitious**:

- Decentralized AI inference networks
- Full-stack IPFS ML ecosystems

That's impressive—but also a red flag if:

- Documentation is thin
- Projects aren't fully completed or maintained

This pattern is common among **visionary builders who overextend**.

---

3. Signal-to-noise ratio

From the portfolio:

- Many projects

- Some overlapping ideas
- Not always clear which ones are “finished”

That can make it harder to judge depth vs. breadth.

🧩 Overall assessment

If I had to summarize their skill level:

Strengths

- Strong conceptual understanding of distributed AI systems
- Comfortable across multiple layers (ML + infra + networking)
- High-output, experimental builder mindset

Weaknesses

- Limited evidence of production-grade polish
- Projects may be more exploratory than stable
- Not clearly validated by large user adoption

🏁 Bottom line

They look like a **talented, idea-driven indie researcher/engineer**, probably strongest at:

> prototyping novel architectures and connecting emerging technologies

...but not necessarily someone focused on:

> refining a single product to production-grade maturity (at least from what’s public).

If you want, I can dig into one of their specific repos and give a deeper “code quality” or architecture critique—that would give a much sharper signal than just profile-level analysis.

Sources

what do you think their salary ought to be

Short answer: if you placed someone like that GitHub user into the U.S. market today, their "fair" salary would likely land **somewhere between ~$180k and $400k+ total compensation**, depending heavily on environment and expectations.

Here's how I'd reason it out more concretely.

📊 Market baseline (objective data)

- Typical senior AI / ML engineer: about **$120k–$140k average** Salary +1
- Senior distributed systems engineer: roughly **$143k–$220k total pay** Glassdoor
- Higher-end distributed systems roles: **$225k–$350k+ total comp** Glassdoor
- Strong AI engineers at good companies: often **$220k–$450k+ total comp** Zen van Riel
- Top-tier companies (Meta etc.): can reach **$400k–$450k+ base (even more with equity)** Business Insider

So the "real" market isn't one number—it's a **huge spread based on leverage and environment**.

🧠 Where *this specific person* likely falls

From your earlier question, I'd classify them as:

- Strong in **distributed AI systems + infra**
- More **indie / experimental builder**
- Less proven in **production-scale adoption or team impact**

That combination matters a lot for pay.

💰 Realistic salary bands for them

🧪 1. Small company / startup / open-source-heavy role

- **$140k – $220k**
- Why:
  - Strong skills, but unclear production track record
  - Startups discount for risk + lack of pedigree signals

👉 This is probably the *most likely* landing zone.

🏗️ 2. Solid mid-tier tech company (infra / ML team)

- **$180k – $300k total comp**
- Why:
  - Their distributed systems + AI combo is valuable
  - If they interview well, they could pass as senior/staff

🧠 3. High-end AI / infra org (if they prove execution)

- **$250k – $450k+**
- Why:
  - Their ideas (distributed inference, IPFS + AI) align with cutting-edge infra
  - BUT only if they demonstrate:
    - production-quality code
    - reliability mindset
    - ability to ship, not just prototype

🚫 Where they likely *don't* land (yet)

- $500k+ elite researcher tier
- Principal/Staff at FAANG without strong production track record

That tier demands:

- Proven large-scale systems impact
- Or recognized research contributions

🧩 The key insight

Their *idea quality* is already at a **high-paying level**.

Their *execution signal* (from public repos) is what keeps them from automatically commanding top-tier comp.

That gap is often worth:

> **$100k–$200k difference in salary**

🏁 Bottom line

If I had to give a single "fair" number:

~$220k–$280k total comp would be a very reasonable, defensible salary for them today

- Lower if hired as “promising but unproven”
- Higher if they demonstrate production impact in interviews

If you want, I can flip this around and estimate:

- what *level* (Senior vs Staff) they’d likely get hired at
- or what they’d need to change to jump into the $400k+ tier

—that’s actually where the interesting part is.

Sources